

w=1024 train-loss trajectories at four learning rates — different ρ values
($\alpha=0.05$, $H=30$ — minimal smoothing to reveal high-LR instabilities)

